





Machine Vision

Lecture 2: Linear Filters


Pinhole camera model

- Light travels on a straight line from the object through then centre of projection
- The image is then created at the intersection with the image plane



Simple lens optics

- In reality lenses have to be used to gather light
- A sharp image is only produced for objects at a specific distance



Simple lens optics

- Unless perfectly focused a blurred image of every scene point is projected onto a circle on the image plane







Image sensing

- The image plane is covered with a grid of **picture cells (aka pixels)** that “count” the number of photons during exposure
- The probability of a photon originating from a scene point being detected by a pixel is proportional to the sensor area covered by its blurred image and the exposure time
- The number of detected photons, and therefore the **intensity value** of each pixel, is affected by quantum fluctuations causing independent **white noise** in each pixel



Linear, shift-invariant systems

- For now we consider an intensity image as 2-dimensional continuous signal

$$f: \mathbb{R}^2 \rightarrow \mathbb{R}$$


- An image processing algorithm is taking images as input and produces a processed image as output
- Let's assume the algorithm would produce the following outputs

$$f_1 \rightarrow g_1$$

$$f_2 \rightarrow g_2$$

- The algorithm is called a **linear system**, if

$$\alpha f_1 + \beta f_2 \rightarrow \alpha g_1 + \beta g_2$$





Linear, shift-invariant systems

- An algorithm is called **shift-invariant**, if the output of a shifted input image is the shifted output image, i.e. if

$$f[x, y] \rightarrow g[x, y]$$

- then

$$f[x - a, y - b] \rightarrow g[x - a, y - b]$$



Linear, shift-invariant systems

- An algorithm that is both linear and shift-invariant is called a **linear, shift-invariant system**
- Every linear, shift-invariant system can be calculated as follows

$$g[x, y] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f[x - \xi, y - \eta] h[\xi, \eta] d\xi d\eta$$

- For some function $h[\xi, \eta]$, called the **point-spread function** or **convolution kernel**
- This operation is called a **convolution**

$$g = f \otimes h$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|--|--|--|
| | | | | |
| | 24 | | | |
| | | | | |
| | | | | |
| | | | | |

$$\begin{aligned} &1 \times 3 \\ &+ 1 \times 3 \\ &+ 2 \times 4 \\ &+ 1 \times 5 \\ &+ 1 \times 5 \\ &= 24 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|--|--|
| | | | | |
| | 24 | 24 | | |
| | | | | |
| | | | | |
| | | | | |

$$\begin{aligned} &1 \times 5 \\ &+ 1 \times 4 \\ &+ 2 \times 5 \\ &+ 1 \times 3 \\ &+ 1 \times 2 \\ &= 24 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | | | | |
| | | | | |
| | | | | |

$$\begin{aligned} &1 \times 3 \\ &+ 1 \times 5 \\ &+ 2 \times 3 \\ &+ 1 \times 5 \\ &+ 1 \times 1 \\ &= 20 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | 21 | | | |
| | | | | |
| | | | | |

$$\begin{aligned} & 1 \times 4 \\ & + 1 \times 3 \\ & + 2 \times 5 \\ & + 1 \times 2 \\ & + 1 \times 2 \\ & = 21 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | 21 | 18 | | |
| | | | | |
| | | | | |

$$\begin{aligned}
 &1 \times 5 \\
 &+ 1 \times 5 \\
 &+ 2 \times 2 \\
 &+ 1 \times 1 \\
 &+ 1 \times 3 \\
 &= 18
 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | 21 | 18 | 14 | |
| | | | | |
| | | | | |

$$\begin{aligned} &1 \times 3 \\ &+ 1 \times 2 \\ &+ 2 \times 1 \\ &+ 1 \times 3 \\ &+ 1 \times 4 \\ &= 14 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | 21 | 18 | 14 | |
| | 20 | | | |
| | | | | |

$$\begin{aligned} &1 \times 5 \\ &+ 1 \times 1 \\ &+ 2 \times 2 \\ &+ 1 \times 3 \\ &+ 1 \times 7 \\ &= 20 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | 21 | 18 | 14 | |
| | 20 | 17 | | |
| | | | | |

$$\begin{aligned} &1 \times 2 \\ &+ 1 \times 2 \\ &+ 2 \times 3 \\ &+ 1 \times 4 \\ &+ 1 \times 3 \\ &= 17 \end{aligned}$$

Convolution

- Intuitively the convolution moves the (flipped) kernel over the input image, multiplies the values and sums them up

| | | | | |
|---|---|---|---|---|
| 1 | 3 | 5 | 3 | 5 |
| 3 | 4 | 5 | 3 | 5 |
| 3 | 5 | 2 | 1 | 3 |
| 1 | 2 | 3 | 4 | 1 |
| 5 | 7 | 3 | 2 | 6 |

$f[x, y]$

\otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

=

| | | | | |
|--|----|----|----|--|
| | | | | |
| | 24 | 24 | 20 | |
| | 21 | 18 | 14 | |
| | 20 | 17 | 15 | |
| | | | | |

$$\begin{aligned} &1 \times 1 \\ &+ 1 \times 3 \\ &+ 2 \times 4 \\ &+ 1 \times 1 \\ &+ 1 \times 2 \\ &= 15 \end{aligned}$$

Point-spread function

- The Dirac δ -function is a generalised function which is zero everywhere, except at the origin where it is “infinite”

$$\delta[x, y] = \begin{cases} \lim_{\epsilon \rightarrow 0} \frac{1}{4\epsilon^2} & \text{if } |x| < \epsilon, |y| < \epsilon \\ 0 & \text{otherwise} \end{cases}$$

- It represents a unit impulse (or a single point) at the origin
- Applying a convolution kernel to this single point yields the kernel itself as result

$$h[x, y] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \delta[x - \xi, y - \eta] h[\xi, \eta] d\xi d\eta$$

Point-spread function

- The convolution of a single point with a kernel yields the kernel as result

$$h = \delta \otimes h$$

- This is the reason it is called the point-spread function, because the kernel determines how a single impulse is spread

| | | | | |
|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 |

$\delta[x, y]$

 \otimes

| | | |
|---|---|---|
| 0 | 1 | 0 |
| 1 | 2 | 1 |
| 0 | 1 | 0 |

$h[-\xi, -\eta]$

 $=$

| | | | | |
|--|---|---|---|--|
| | | | | |
| | 0 | 1 | 0 | |
| | 1 | 2 | 1 | |
| | 0 | 1 | 0 | |
| | | | | |


$h[x, y]$

Blurring

- Applying a kernel with only positive numbers spreads out intensity
- These kernels are the basis for all blurring operations



```
kernel = np.array([[1,1,1,1,1],  
                    [1,1,1,1,1],  
                    [1,1,1,1,1],  
                    [1,1,1,1,1],  
                    [1,1,1,1,1]])/25  
result = cv2.filter2D(img, -1, kernel)
```





Edge detection

- Edges are local image features where brightness changes between two areas
- Edges in images typically derive from occluding contours, i.e. where one object is in front of another object
- A brightness edge is locally

$$E[x, y] = B_1 + (B_2 - B_1)u[x \sin \theta + y \cos \theta + \rho]$$

- Using the step function

$$u[z] = \int_{-\infty}^z \delta[t] dt$$





Edge detection

- Edges are local image features where brightness changes between two areas
- Edges in images typically derive from occluding contours, i.e. where one object is in front of another object
- The derivatives are only non-zero on the edge and proportional to the contrast of the edge

$$\frac{\partial E}{\partial x} = \sin \theta (B_2 - B_1) \delta[x \sin \theta + y \cos \theta + \rho]$$

$$\frac{\partial E}{\partial y} = -\cos \theta (B_2 - B_1) \delta[x \sin \theta + y \cos \theta + \rho]$$



Edge detection

- Edges are local image features where brightness changes between two areas
- Edges in images typically derive from occluding contours, i.e. where one object is in front of another object
- The derivatives are only non-zero on the edge and proportional to the contrast of the edge
- Therefore, to find edges in the image (i.e. the object boundaries) it is useful to look at the image gradient





Image gradients

- The derivative is a linear, shift-invariant operation, therefore it can be computed using convolution

$$\frac{\partial E}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{E[x + \epsilon, y] - E[x - \epsilon, y]}{2\epsilon} = E \otimes \delta_x[x, y]$$

$$\frac{\partial E}{\partial y} = \lim_{\epsilon \rightarrow 0} \frac{E[x, y + \epsilon] - E[x, y - \epsilon]}{2\epsilon} = E \otimes \delta_y[x, y]$$

with

$$\delta_x[x, y] = \lim_{\epsilon \rightarrow 0} \frac{\delta[x + \epsilon, y] - \delta[x - \epsilon, y]}{2\epsilon}$$

$$\delta_y[x, y] = \lim_{\epsilon \rightarrow 0} \frac{\delta[x, y + \epsilon] - \delta[x, y - \epsilon]}{2\epsilon}$$

Edge detection



- The discrete version of the two doublets

$$\delta_x[x, y] = \lim_{\epsilon \rightarrow 0} \frac{\delta[x + \epsilon, y] - \delta[x - \epsilon, y]}{2\epsilon}$$

$$\delta_y[x, y] = \lim_{\epsilon \rightarrow 0} \frac{\delta[x, y + \epsilon] - \delta[x, y - \epsilon]}{2\epsilon}$$

- Allows to calculate derivatives of images using convolution filters

```
dx = np.array([[1, 0, -1]])  
dy = np.array([[1],  
               [0],  
               [-1]])  
  
fx = cv2.filter2D(img, -1, dx)  
fy = cv2.filter2D(img, -1, dy)
```



Properties of convolutions

- Convolutions are commutative, i.e.

$$g = f \otimes h = h \otimes f$$

- and associative

$$g = (f \otimes h_1) \otimes h_2 = f \otimes (h_1 \otimes h_2)$$

- Therefore, the order in which convolution operations are applied does not matter
- In particular, convolutions can be pooled together and applied as one single convolution to an input image

The modulation-transfer function

- Applying a convolution to the complex function

$$f[x, y] = e^{i(ux+vy)} = \cos[ux + vy] + i \sin[ux + vy]$$

- Yields

$$g[x, y] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} e^{i(u(x-\xi)+v(y-\eta))} h[\xi, \eta] d\xi d\eta$$

$$= e^{i(ux+vy)} \underbrace{\int_{-\infty}^{\infty} \int_{-\infty}^{\infty} e^{-i(u\xi+v\eta)} h[\xi, \eta] d\xi d\eta}_{H(u,v)}$$

- The **modulation transfer function** $H(u, v)$ does not depend on x or y , therefore the convolution operation has a multiplicative effect on individual frequencies (i.e. $e^{i(ux+vy)}$ is an eigenfunction of convolution in two dimensions)

Fourier transformation

- Every function can be considered as a sum of an infinite number of sinusoidal waves

$$f[x, y] = \frac{1}{4\pi^2} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F[u, v] e^{i(ux+vy)} du dv$$

- This frequency representation $F[u, v]$ is called the **Fourier transformation** of the function $f[x, y]$

- Using the modulation transfer function the convolution is then





$$g[x, y] = f \otimes h = \frac{1}{4\pi^2} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} H[u, v] F[u, v] e^{i(ux+vy)} du dv$$

- Now looking at the Fourier transformation of $g[x, y]$ one can see that

$$G[u, v] = F[u, v] H[u, v]$$

- **Convolution is simply a multiplication in the frequency domain!**


Fourier spectrum



Fourier spectrum



$f[x, y]$



$|F[u, v]|$

Machine Vision

Convolution in the frequency domain

- Convolution in the spatial domain

$$g[x, y] = f[x, y] \otimes h[x, y]$$

- Becomes multiplication in the frequency domain

$$G[u, v] = F[u, v]H[u, v]$$

- The modulation transfer function attenuates/dampens certain frequencies and directions
- White noise is uniform across the spectrum, so the modulation transfer function also tells us what noise components are attenuated and what noise components are dampened by the application of a filter
- We can also use this to understand how convolution filters can be inverted by observing that this is only possible if the kernel spectrum is non-zero everywhere, in which case the inverse is simply


$$F[u, v] = \frac{G[u, v]}{H[u, v]}$$

Linear Filters and FFT

- To calculate the convolution

$$g = f \otimes x$$

- We can do the following
 - Calculate Fourier transform of f and h
 - Calculate $G = FH$ by multiplication
 - Calculate the inverse Fourier transform of G
- Because spatial convolution requires $O(n^2)$ multiplications and the Fast-Fourier-Transform algorithm runs in $O(n \log n)$ this strategy can be faster (usually depending on the size of h)





FFT in Python

```
sigma = 50

x,y = np.meshgrid(np.arange(0,len(img[0])),np.arange(0,len(img)))
kernel = np.exp(-((x-len(img[0])/2)**2+(y-len(img)/2)**2)/(2*sigma**2))/(2*np.pi*sigma**2)

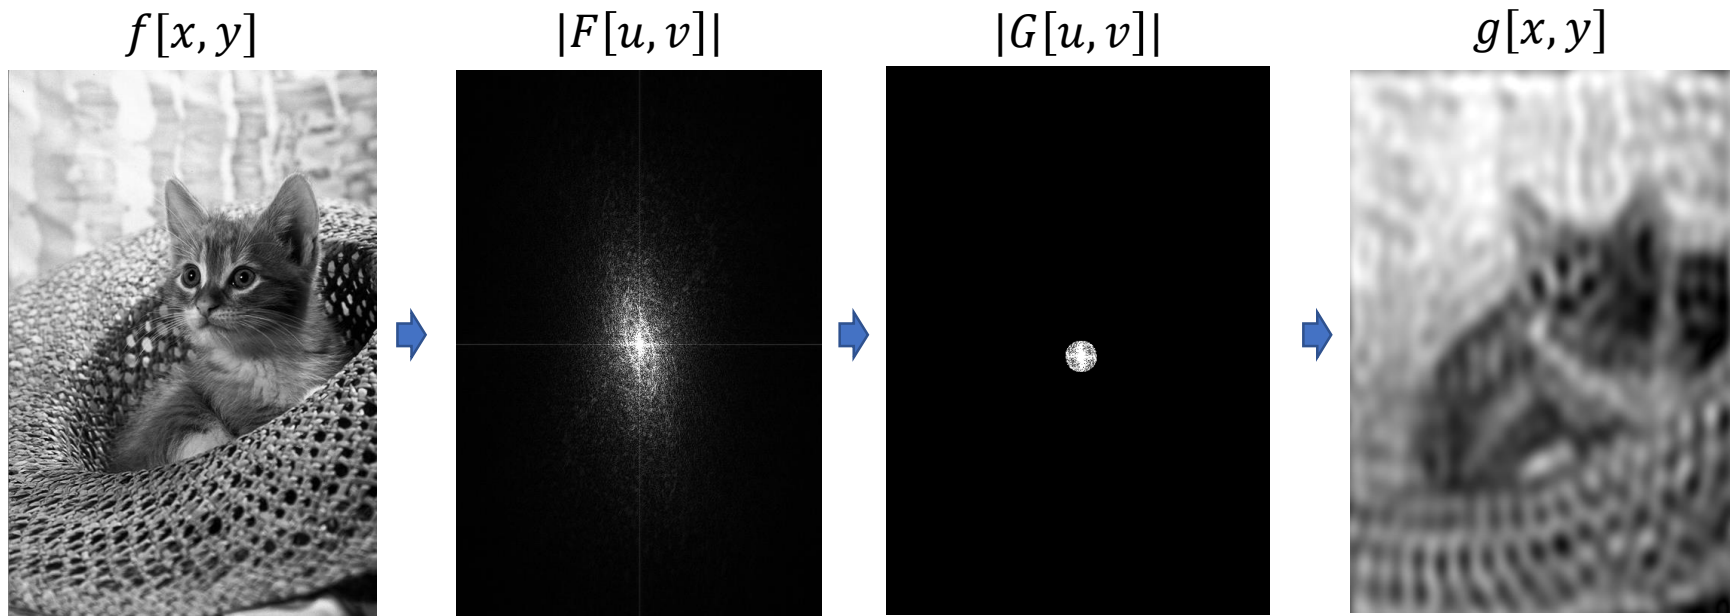
ft_img = np.fft.fft2(img)
ft_kernel = np.fft.fft2(np.fft.fftshift(kernel))

result = abs(np.fft.ifft2(ft_img * ft_kernel))/255
# result = cv2.filter2D(img, -1, kernel)
cv2.imshow("result", result)
```




Low-pass filter

- Removing high frequencies from the image is called **low-pass** filtering
- It only retains homogeneous areas and removes all high-frequency edges
- The result is a blurred image



High-pass filter

- Removing low frequencies from the image is called **high-pass** filtering
- It only retains edges and removes all low-frequency homogeneous areas
- The result is an image with homogeneous areas removed



Smoothing at different scales


- We can filter the image with Gaussians of different size

$$g_{\sigma} = \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2}$$


- To generate images at different smoothing scales

$$\bar{f} = g_{\sigma} \otimes f$$


$\sigma = 10$







$\sigma = 20$




$\sigma = 30$



$\sigma = 40$



Machine



Derivatives at different scales

- Calculating derivatives on a discrete image using convolution filters is not practical (because it would require infinitely small pixels)

$$\delta_x[x, y] = \lim_{\epsilon \rightarrow 0} \frac{\delta[x + \epsilon, y] - \delta[x - \epsilon, y]}{2\epsilon}$$

$$\delta_y[x, y] = \lim_{\epsilon \rightarrow 0} \frac{\delta[x, y + \epsilon] - \delta[x, y - \epsilon]}{2\epsilon}$$

- Instead, we have to calculate derivatives on a smoothed version of the image (remember, convolution is associative and commutative)

$$\bar{f}_x = (\delta_x \otimes g_\sigma) \otimes f$$


$$\bar{f}_y = (\delta_y \otimes g_\sigma) \otimes f$$

Derivatives at different scales

- We can calculate the derivative of the Gaussian function directly

$$\frac{\partial g_{\sigma}}{\partial x} = -\frac{x}{2\pi\sigma^4} \exp -\frac{x^2 + y^2}{2\sigma^2}$$

$$\frac{\partial g_{\sigma}}{\partial y} = -\frac{y}{2\pi\sigma^4} \exp -\frac{x^2 + y^2}{2\sigma^2}$$



- And calculate the (scale-dependent) derivatives using the following convolutions

$$\bar{f}_x = \frac{\partial g_{\sigma}}{\partial x} \otimes f$$

$$\bar{f}_y = \frac{\partial g_{\sigma}}{\partial y} \otimes f$$

Edge detection

- If we want to have a direction independent edge detector we can look at the squared gradient image

$$\left(\frac{\partial g_\sigma}{\partial x} \otimes f \right)^2 + \left(\frac{\partial g_\sigma}{\partial y} \otimes f \right)^2$$

- The Derivative of Gaussian operator is still scale-dependent, which suggests that every edge does not only have a particular direction, but also a particular scale
- We will see later how scale invariant feature detection can be achieved (SIFT)

$\sigma = 1$



$\sigma = 2$




$\sigma = 3$




Machine Vision

$\sigma = 4$



$\sigma = 5$





Sobel operator

- In practical applications it makes sense to choose convolution kernels that approximate derivatives of Gaussians while at the same time are efficient to compute
- The **Sobel operator** applies a 3x3 convolution filter which applies a difference operation in x-direction and a smoothing operation in y-direction (and vice versa) similar to a derivative of Gaussian operator
- The scale in this case is fixed to the image resolution

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

| | | |
|----|----|----|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| -1 | -2 | -1 |



Summary

- Every linear, shift-invariant system can be expressed as convolution
- A convolution in the spatial domain corresponds to a multiplication in the frequency domain
- Image smoothing and image derivatives are two important examples of linear, shift-invariant operations
- Occluding object boundaries create edges in images, which can be detected using such linear, shift-invariant filters
- These edge features are characterised by their direction and by their scale

Thank you for your attention!