

Machine Learning



Machine Learning

Lecture: Instance-Based Learning

Ted Scully

k - Nearest Neighbour Distance Metric

- ▶ There are a number of distance measurements used to determine similarity:
 - ▶ Euclidean Distance
 - ▶ Manhattan Distance
 - ▶ Minkowski
 - ▶ Hamming Distance

Distance Metrics

- ▶ An important aspect of k-NN algorithms is how we determine which instances are the nearest to the target case. Thus, the distance metric is a measure of the similarity between two cases.
- ▶ Common distance metrics include:
 - ▶ Euclidean
 - ▶ Manhattan
 - ▶ Minkowski
- ▶ To help illustrate the various metrics let's assume we have the dataset below with n features and two instances p and q

	Feature 1	Features 2	Feature n
p	p_1	p_2	p_n
q	q_1	q_2	q_n

Euclidean Distance Metric

- ▶ The most common measure of distance is **Euclidian distance**: which measures the straight-line distance between two points.
- ▶ If $\mathbf{p} = \langle p_1, p_2, \dots, p_n \rangle$ and $\mathbf{q} = \langle q_1, q_2, \dots, q_n \rangle$ are two points in Euclidean n-space, then the distance (d) from p to q, or from q to p is given by

$$\begin{aligned} d(\mathbf{p}, \mathbf{q}) &= d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \dots + (q_n - p_n)^2} \\ &= \sqrt{\sum_{i=1}^n (q_i - p_i)^2}. \end{aligned}$$

It is important to understand that p and q here represent two data instances. Each instance consisting of a finite set of features. The instance q has the features q1, q2, ... qn.

Euclidean Distance Metric

Movie title	# of clicks	# of classes	Movie title	Distance to movie “?”
California Man	3	104	California Man	20.5
He’s Not Really into Dudes	2	100	He’s Not Really into Dudes	18.7
Beautiful Woman	1	81	Beautiful Woman	19.2
Kevin Longblade	101	10	Kevin Longblade	115.3
Robo Slayer 3000	99	5	Robo Slayer 3000	117.4
Amped II	98	2	Amped II	118.9
?	18	90	Action	
			Unknown	

Distance between **California man** (3, 104) and the **query** instance (18, 90) would be:

Euclidean Distance Metric

Movie title	# of clicks	# of classes	Movie title	Distance to movie “?”
California Man	3	104	California Man	20.5
He’s Not Really into Dudes	2	100	He’s Not Really into Dudes	18.7
Beautiful Woman	1	81	Beautiful Woman	19.2
Kevin Longblade	101	10	Kevin Longblade	115.3
Robo Slayer 3000	99	5	Robo Slayer 3000	117.4
Amped II	98	2	Amped II	118.9
?	18	90	Action	
			Unknown	

Distance between **California man** (3, 104) and the **query** instance (18, 90) would be


$$\sqrt{(3 - 18)^2 + (104 - 90)^2} = \sqrt{225+196}= 20.5$$

Manhattan Distance Metric


- ▶ Manhattan distance measures distance **parallel to each axis**, not diagonally (in downtown Manhattan, to get from one point to another you generally walk North-South and East-West, rather than 'as the crow flies').
- ▶ In other words take the sum of the absolute values of the differences of the coordinates
- ▶ $d(p, q) = |q_1 - p_1| + |q_2 - p_2| + |q_n - p_n| = \sum_{i=1}^n |q_i - p_i|$

Manhattan Distance Metric

- ▶ Lets assume we have a simple dataset containing three instances as follows. We are also given the query instance below. Calculate the Manhattan and Euclidean distance between the **first training example** and the **query**



Manhattan



Euclidean

Area	Weight	Height	Capacity
10	8	4	14
12	10	6	12
14	9	4	11

Area	Weight	Height	Capacity
5	4	4	10

Area	Weight	Height	Capacity
10	8	4	8
12	10	6	12
14	9	4	11

Area	Weight	Height	Capacity
5	4	4	3

Manhattan Distance Metric

- ▶ Euclidean

- ▶ $(10-5)^2 + (8-4)^2 + (4-4)^2 + (8-3)^2$

- ▶ 57

- ▶ Square root of 57 = 7.55

- ▶ Manhattan

- ▶ $|10-5| + |8-4| + |4-4| + |8-3|$

- ▶ $5 + 4 + 0 + 5 = 14$

Area	Weight	Height	Capacity
10	8	4	8
12	10	6	12
14	9	4	11

Area	Weight	Height	Capacity
5	4	4	3

Minkowski Distance

- ▶ The Minkowski distance between a feature vector $\mathbf{p} = \langle p_1, p_2, \dots, p_n \rangle$ and another feature vector $\mathbf{q} = \langle q_1, q_2, \dots, q_n \rangle$ is defined as:
- ▶
$$d(p, q) = \left(\sum_{i=1}^n |p_i - q_i|^a \right)^{\frac{1}{a}}$$
- ▶ In the above equation a is an integer. Consider the case when $a = 1$ or $a = 2$. Any comments.

Minkowski Distance

- ▶ The Minkowski distance between a feature vector $\mathbf{p} = \langle p_1, p_2, \dots, p_n \rangle$ and another feature vector $\mathbf{q} = \langle q_1, q_2, \dots, q_n \rangle$ is defined as:
- ▶
$$d(p, q) = \left(\sum_{i=1}^n |p_i - q_i|^a \right)^{\frac{1}{a}}$$
- ▶ In the above equation a is an integer. Consider the case when $a = 1$ or $a = 2$. Any comments.
 - ▶ For $a = 1$ we get the Manhattan distance and for $a = 2$ we get the Euclidean distance.
 - ▶ Minkowski Distance is a generalization of the Euclidean and Manhattan distance metrics.

Hamming Distance Metric

- ▶ Allows us to deal with problems that have features that are categorical (discrete) rather than **continuous-valued**.
- ▶ The value 0 is assigned for each feature where both cases have the same value, 1 for each where they are different.

- ▶
$$d(p_i, q_i) = \begin{cases} 0 & \text{if } q_i == p_i \\ 1 & \text{if } q_i \neq p_i \end{cases}$$

ID	Outlook	Temp	Hum	Windy	Play?
A	sunny	hot	high	false	no
B	sunny	hot	high	true	no
C	overcast	hot	high	false	yes

What is distance
between B and C?

Outlook (1) + Temp(0) +
Hum(0) + Windy(1) = 2

Heterogeneous Distance Metric

- ▶ When we have a dataset that is a mix of discrete and continuous valued features we can combine distance measures such as Manhattan and Hamming.

$$d(p_i, q_i) = \begin{cases} |q_i - p_i| & \text{if feature } i \text{ is continuous} \\ 0 & \text{if feature } i \text{ is discrete and } q_i == p_i \\ 1 & \text{if feature } i \text{ is discrete } q_i \neq p_i \end{cases}$$

The Hamming metric is limited and provides limited information about the difference between features.

Building a K Nearest Neighbour Classifier

- ▶ **Step 1. Read information from a dataset**
 - ▶ Read data from a dataset containing classified instances. Read each feature of the dataset as well as corresponding class.
- ▶ **Step 2. Determine distance between each dataset entry and the query instance**
 - ▶ Use a suitable distance metric to calculate the distance between the query instance and all k neighbours
- ▶ **Step 3. Classify the query instance**
 - ▶ Identify k nearest data instances. Assign query instance category corresponding to most common category.

Problems Measuring Distance 1 -Scale

- ▶ Since the performance of k-NN is strongly dependent on the choice of distance metric, you need to be aware of some pitfalls.
- ▶ The first problem arises when the features are different from each other.
- ▶ For example, if one feature has a range between **0 and 1** and another feature has a range between **0 and 10, 000**, it hardly makes sense to add them as would happen with Euclidian or Manhattan distance metrics (for example, salary and age).
- ▶ What is the main problem that arises from the above situation?

Problems Measuring Distance (1)

- ▶ **Problem 1: Scaling**

- ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Feature B will dominate calculations

- ▶ Example, lets calculate the distance between data instance 1 and 2
 - ▶ Data instance 1 = (5.5, 787)
 - ▶ Data instance 2 = (7.5, 567)

Problems Measuring Distance (1)

- ▶ **Problem 1: Scaling**

- ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Feature B will dominate calculations

- ▶ Example, lets calculate the distance between data instance 1 and 2
 - ▶ Data instance 1 = (5.5, 787)
 - ▶ Data instance 2 = (7.5, 567)

$$\sqrt{(5.5 - 7.5)^2 + (787 - 567)^2}$$

$$\sqrt{16 + 48400}$$

Problems Measuring Distance (1)

▶ Problem 1: Scaling

- ▶ Feature A has range 1-10
Feature B has range 1-1000
- ▶ Feature B will dominate calculation

We can see below that the second feature is entirely dominating the distance calculation simply because it has a larger range of values compared to the first feature.

▶ Example, lets calculate the distance

- ▶ Data instance 1 = (5.5, 787)
- ▶ Data instance 2 = (7.5, 567)

We don't want our model to bias toward a particular feature simply because the range happens to be larger.

$$\sqrt{(5.5 - 7.5)^2 + (787 - 567)^2}$$

$$\sqrt{16 + 48400}$$

Problems Measuring Distance (1)

- ▶ Solution:
 - ▶ Normalise all dimensions independently (scale data so that it has a maximum and minimum range)
 - ▶ Using range normalization we identify the minimum and maximum value for a specific feature. We can then apply the following formula.
 - ▶
$$\text{newValue} = \frac{\text{originalValue} - \text{minValue}}{\text{maxValue} - \text{minValue}}$$

$$\triangleright \text{newValue} = \frac{\text{originalValue} - \text{minValue}}{\text{maxValue} - \text{minValue}}$$

▶ Problem 1: Scaling

- ▶ Feature A has range 1-10
Feature B has range 1-1000

▶ Normalise variables

- ▶ Feature A
 - ▶ $(5.5 - 1)/(10-1) = 0.5$
 - ▶ $(7.5 - 1)/(10 - 1) = 0.72$
- ▶ Feature B
 - ▶ $(787-1)/(1000-1) = 0.78$
 - ▶ $(567-1)/(1000-1) = 0.56$

▶ Before Normalization

- ▶ Data instance 1 = (5.5, 787)
- ▶ Data instance 2 = (7.5, 567)

▶ After Normalization

- ▶ Data instance 1 = (0.5, .78)
- ▶ Data instance 2 = (0.72, 0.56)

$$\sqrt{(0.5 - 0.72)^2 + (0.78 - 0.56)^2}$$

$$\sqrt{0.048 + 0.048}$$

Problems Measuring Distance (1)

- ▶ When we normalize the train data, it is also important to understand:
 - ▶ We normalize each feature **independently**
 - ▶ We must normalize the test data using the same parameters for max and min (that is we still use the minValue and maxValue from the original training set).

Problems Measuring Distance – Irrelevant Features

- ▶ The other principal problem is that **all features are included equally** in the calculations we have looked at, even though some features may be **redundant** or **less relevant**.
- ▶ Therefore, a number of features may skew the result even though they might have little or no impact on the classification.
- ▶ **Solution 2A:**
 - ▶ **Assign weighting** to each dimension (Optimise weighting to minimise error)
- ▶ **Solution 2B:**
 - ▶ Give some dimensions **0 weight** (Feature subset solution)
- ▶ Either way, since we cannot know in advance what weighting to give dimensions, systematic repeated experiments are needed to optimise them.
- ▶ We will look at feature selection in more detail later in the module.

Problem with KNN


Assume in the following example that we are using $k=7$?

What are the potential problems?

What might you do to address this problem?

Distance-Weighted kNN

Give each neighbour weight: inverse of distance from target



Distance Weighted k-NN (Regression)

- It is the vote of each neighbour that is weighted according to how close it is to the target, so the **closer neighbour's influence the prediction more**.
- We can if we wish use all cases as neighbours, since those very far from the target will have little influence on the prediction, but none will be completely ignored.
- We do the following for a regression problem.

Given a query instance x_q ,

$$f(\mathbf{x}_q) := \frac{\sum_{i=1}^k w_i f(\mathbf{x}_i)}{\sum_{i=1}^k w_i}$$

Where

$$w_i = \frac{1}{d(\mathbf{x}_q, \mathbf{x}_i)^2}$$

Distance Weighted k-NN (Regression)

- It is the vote of each neighbour that is weighted according to how close it is to the target, so the **closer neighbour's influence the prediction more**.
- We can if we wish use all cases as neighbours, since those very far from the target will have little influence on the prediction, but none will be completely ignored.
- We do the following for a regression problem.

Given a query instance x_q ,

$$f(\mathbf{x}_q) := \frac{\sum_{i=1}^k w_i f(\mathbf{x}_i)}{\sum_{i=1}^k w_i}$$

Where

$$w_i = \frac{1}{d(\mathbf{x}_q, \mathbf{x}_i)^2}$$


Here you will notice that we use the inverse distance to the power of 2 ($n=2$). This is typical. However, we can also use $n=1$, $n=3$, etc

Distance Weighted k-NN Algorithm

Regression Example

- ▶ Consider the basic regression example depicted in the slide.
- ▶ Distance between query instance and:
 - ▶ Case 82 is 2
 - ▶ Case 83 is 4
 - ▶ Case 78 is 6
- ▶ What is the value of the query instance.


$$f(\mathbf{x}_q) := \frac{\sum_{i=1}^k w_i f(\mathbf{x}_i)}{\sum_{i=1}^k w_i}$$



Distance Weighted k-NN Algorithm

Regression Example


- ▶ Distance between query instance and:
 - ▶ Case 82 is 2
 - ▶ Case 83 is 4
 - ▶ Case 78 is 6



Distance Weighted k-NN Algorithm

Regression Example

- ▶ Distance between query instance and:
 - ▶ Case 82 is 2
 - ▶ Case 83 is 4
 - ▶ Case 78 is 6



- ▶ $((1/4)(82) + (1/16)(83) + (1/36)(78)) / (1/4 + 1/16 + 1/36)$
- ▶ $= 27.854 / 0.34027777$
- ▶ $= 81.856$

Distance Weighted k-NN (Classification)

- We iterate through each class. For a specific class we identify each instances amongst the k nearest instances that belong to that class. We then add up the inverse distance for each of the identified instances.
- The class that results in the largest value is the selected class for the new query instance.

$$vote(c_j) := \sum_{i=1}^k \frac{1}{d(\mathbf{x}_q, \mathbf{x}_i)^n} (c_i, c_j)$$

(y_i, y_j) returns 1 if the class labels match and 0 otherwise

What do you think will be the impact of n (n must be a positive number greater than or equal to 1)

Vote(Purple Class) (n=1)

Purple

1/10 +
1/9 +
1/5 = 0.41

Yellow

1/5 +
1/6 +
1/5 = 0.566

Black

1/1 +
1/2 +
1/2 = 2

Purple


10
9
5

Yellow

5
6
5

Black

1
2
2



$$vote(y_j) := \sum_{i=1}^k \frac{1}{d(\mathbf{x}_q, \mathbf{x}_i)^n} (y_i, y_j)$$

Result of Voting

- ▶ Result of voting is
- ▶ $V(\text{Purple}) = 0.41$
- ▶ $V(\text{Yellow}) = 0.566$
- ▶ $V(\text{Black}) = 2$
- ▶ Therefore the query instance is classified as a **Black class**.

Assessing the Performance of a Regression Model

- ▶ So far we have a basic measure that we can use for assessing the performance of a classification model, which is the **accuracy**.
- ▶ $\text{Accuracy} = \frac{\text{number of test instances correctly classified}}{\text{total number of test instances}}$.
- ▶ Later in the module we will look more comprehensively at evaluation metrics.
- ▶ So what is a common **evaluation metric** we can use for **regression**?

Basic Measures of Error (Regression)


- The R^2 coefficient compares the **performance of a model on a test set (sum of squared residuals)** with the performance of an imaginary model that always predicts the **average values from the test set (total sum of squares)**.
- The R^2 coefficient is calculated as:

$$R^2 = 1 - \frac{\text{sum of squared residuals}}{\text{total sum of squares}}$$

- Where

$$\text{sum of squared residuals} = \sum_{i=0}^m (f(x^i) - y^i)^2$$

$$\text{total sum of squares} = \sum_{i=0}^m (\bar{y} - y^i)^2$$



Basic Measures of Error (Regression)

- The R^2 coefficient values typically fall in the range $[0, 1)$ and larger values indicate better model performance.
- The **worse the model** produced, the closer the sum of square residuals value will be to the total sum of squares value. Consequently the **smaller the total R^2** .
- The **better the model** the smaller the squared residuals (smaller error in the model) and the **larger the R^2** value.
- While it is rare, the model produced could be worse than the total sum of squares. In this case the R^2 would be **negative**. The worse the model the lower the R^2 values. It means that whatever model that you came up with is worse than predicting the mean (not a good sign!).

Eager vs. Lazy Learner (1)

- ▶ Eager Learning (Such as Bayesian classifiers, decision trees, neural networks)
 - ▶ When given training data, it **constructs a model** for future use in prediction that summarises the data
 - ▶ **Slow in model construction**, generally quick when classifying unseen instances
- ▶ Instance based learning often referred to as lazy Learners
 - ▶ **No** explicit global **model** constructed
 - ▶ **Calculations deferred** until new case to be classified
 - ▶ Creates **many local approximations**, whereas eager learners create a global approximation
 - ▶ **Significant calculations** needed to take place for each new query (can be slow)

When to use k-Nearest Neighbour

- Primary Benefits
 - **Comprehensibility**: easy to understand
 - Relatively straight-forward to **implement**
 - Can easily handle **multi-class** datasets
 - Effective classifiers for **complex target functions** => good for diverse concepts.
 - Can be used for both **regression and classification** problems and can mix **feature types** in the one dataset so they are very flexible.
- Consider using when:
 - Moderate number of **training instances**
 - Moderate number of **features** per instance (< 20) [Note: If dealing with datasets with a large number of features we can perform dimensionality reduction and still use a k-NN approach]